

CFI, SFI, and All That Jazz

Benjamin Lim
@jarsp






“I want to run **untrusted** (but not actively malicious) code...”





“I want to run **untrusted** (but not actively malicious) code...

...but **VMs/containers** are **insufficient/unsupported**”



Why?




... You want to **run code in the kernel**

... You are on an **embedded system** without the appropriate support

... You are a Cyber Reasoning System and the **humans** didn't give you any other choice

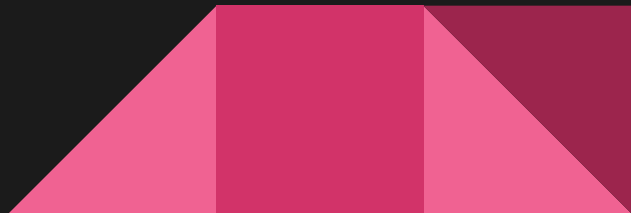
... You want it as a **defense-in-depth** measure



What we'll be covering



- ❖ Overview of techniques to protect and sandbox code, namely:
- ❖ **Software Fault Isolation (SFI)**, *or*, how to protect running code from other code
- ❖ **Control Flow Integrity (CFI)**, *or*, how to protect running code from itself



Software Fault Isolation



- ❖ Prevent untrusted code from escaping sandbox
- ❖ Prevent untrusted code from tampering with trusted components

Portable Native Client (PNaCl)




- ❖ Google's implementation of SFI (IL -> asm)
 - x86, AMD64, ARM
- ❖ WebAssembly precursor
- ❖ Somewhere between **-9% (!) to big% overhead**
- ❖ Performance hit mitigated by speculative execution
- ❖ Implemented with modified GCC for AMD64

Portable Native Client (PNaCl)



- ❖ PNaCl setup:
 - Restrict to 4GB of addressable memory (ARM32 source compatibility)
 - Reserved register (r15) to reference start of untrusted address space
 - Unmapping 10 x 4GB of memory on both sides of untrusted address space
- ❖ Relies on page fault mechanism and zeroing behavior of 32-bit arithmetic



Portable Native Client (PNaCl)



```
mov %rax, %rsp
```

```
mov %eax, %esp
```

```
lea %(r15,%rsp,1), %rsp
```

```
add $0x8, %rcx
```

```
add $0x8, %ecx
```

```
mov %eax, $disp(,%rcx,scale)
```

```
mov %eax, $disp(%r15,%ecx,scale)
```


```
mov $disp(%rsp), %eax
```

```
mov $disp(%rsp), %eax
```

Are we done?



no



Portable Native Client (PNaCl)



- ❖ Statically link code, must be read-only
- ❖ **Bundle and align** valid instructions to 32 bytes
- ❖ **Indirection only to start of bundle**

```
jmp %rdx  
    ↓  
and $0xfffffffffe0, %edx  
lea (%r15,%edx,1), %rdx  
jmp *%rdx
```

Verified SFI



- ❖ Similar work (Kroll et. al. '14)
- ❖ **Formally specify and guarantee SFI** properties of generated code
 - Memory safety
 - Ideally provable functional correctness
 - **Pain and Suffering**
- ❖ Proof using a proof assistant (F^* , Coq) to produce a certificate/extracted program

Typical Exploit A



- ❖ Application accepts user input without proper validation
- ❖ Corruption of critical structures eventually results in code execution
- ❖ Code execution requires control over control flow
- ❖ Targets of interest:
 - Indirect branches and calls
 - Saved return addresses

Control Flow Integrity



- ❖ Force control along 'benign' routes
- ❖ Various granularities possible, leading to differing amounts of overhead
- ❖ Static analysis to recover CFG
- ❖ Compile-time/binary instrumentation to enforce

LLVM CFI



❖ Lightweight forward CFI

- Virtual Calls
- Indirect Function Calls

❖ Backwards CFI planned

- Return Elision (?) for leaf functions
- Explicit call-site checking
- etc.


❖ Shadow Stack option

```
ca7fbb:      mov     (%rdi),%rcx
ca7fbe:      lea     0x7fb42c3(%rip),%rdx
ca7fc5:      mov     %rcx,%rax
ca7fc8:      sub     %rdx,%rax
ca7fcb:      rol     $0x3d,%rax
ca7fcf:      cmp     $0x17f,%rax
ca7fd5:      ja      ca8511
ca7fdb:      lea     0x6f70bc0(%rip),%rdx
ca7fe2:      testb   $0x10, (%rax,%rdx,1)
ca7fe6:      je      ca8511
ca7fec:      callq    *0x98(%rcx)
[... ]
ca8511:      ud2
```

MITRE Embedded CTF 2019



- ❖ Implement a simple 'game console'
- ❖ (Presumably) Vulnerable C/C++ game binaries
- ❖ Assume some basic 'sanity' of binaries
- ❖ Prevent plaintext dump of game
- ❖ No criteria on game performance!




Shadow Stack




- ❖ Calls should generally be paired with returns
- ❖ Maintain a separate 'shadow' stack
- ❖ On return, abort if addresses do not match up
- ❖ Intel Control-Flow Enforcement Technology (CET) proposal

Shadow Stack



'Classic' Shadow Stack

Shadow Stack



'Sparse' Shadow Stack

Shadow Stack Implementation




❖ Dynamic instrumentation framework

- Intel PinTools: x86, AMD64
- DynamoRIO: x86, AMD64, ARM

❖ ARM-to-ARM JIT

❖ Hooks

- Basic block creation
- Signals
- Syscalls
- ...



Shadow Stack Implementation



❖ Basic idea:

- Maintain a (thread-local) **array of return addresses**
- Instrument call and return instructions to **push and pop return addresses**

```
if (tls->shadow_count >= MAX_SHADOW_SIZE)
{
    DERR("Call stack depth exceeded.\n");
    dr_abort();
}
```

```
res = (next_h << 16) + next_l;
tls->shadow_stack[tls->shadow_count].lr = res;
tls->shadow_count++;
```

Shadow Stack Implementation



- ❖ What is a call and what is a return?
 - ARM calls: `bl`, `blx`
 - ARM returns: `bx lr`, `(ldr pc, [sp], #4)`, `ldmia {...,pc}`
- ❖ When to instrument?
 - **Before calls** (at start of function is possible but harder)
 - **Before returns** (to catch bad returns)

Shadow Stack Implementation



❖ Dynamorio/ARM specific issues

- DynamoRIO thinking that **PLT dispatch** is return
 - Special case to eliminate this
- **ldrex/strex**
 - Usually not a problem
- **ite**
 - Insert instrumentation before **ite** block
 - Should emulate

Shadow Stack Implementation



❖ Dynamorio/ARM specific issues

- **ite/predicated returns**
 - Instrumentation checks if return occurs

```
dr_get_mcontext(drcontext, &mc);  
cpsr = mc.cpsr;  
N = ((cpsr & 0x80000000) != 0);  
Z = ((cpsr & 0x40000000) != 0);  
C = ((cpsr & 0x20000000) != 0);  
V = ((cpsr & 0x10000000) != 0);  
switch (exit_pred) {  
    case DR_PRED_EQ:  
        cond = Z;  
        break;  
    ...  
    if (!cond) return;
```


Shadow Stack Implementation



❖ Program loading breaks

- Extract address of main from CRT code
- Instrument but don't execute until main reached

```
10448:      f8df c010      ldr.w   ip, [pc, #16]    ; 1045c <_start+0x24>
1044c:      f84d cd04      str.w   ip, [sp, #-4]!
10450:      4803          ldr     r0, [pc, #12]    ; (10460 <_start+0x28>)
10452:      4b04          ldr     r3, [pc, #16]    ; (10464 <_start+0x2c>)
10454:      f01a f87a      bl      2a54c <__libc_start_main>
10458:      f01f f946      bl      2f6e8 <abort>
```

Shadow Stack Implementation



❖ Intentional Call/Ret mismatch

- setjmp/longjmp
- try/catch (C++)

❖ Solution:

- Heuristically identify key functions
 - __sigsetjmp, __longjmp
 - __restore_core_regs, something else I forgot
- Keep sp along with lr
- Persist entry until function it is called from returns
- Unroll until matching sp address

Conclusion



Cool



End.



Questions?

